



AWS
re:Invent

CON306

Building machine-learning infrastructure on Amazon EKS with Kubeflow

Yaniv Donenfeld

Sr. BDM, Container Services
AWS

Jean Marie Ferdegue

Director, Platform
Engineering, Babylon Health

Jeremie Vallee

AI Infrastructure Lead
Babylon Health

Agenda

AWS

- The AWS ML stack
- Machine learning - Why Kubernetes?
- Kubeflow and Kubeflow pipelines
- Common requirements
- Making AWS a first-class citizen of Kubeflow

Babylon Health

- The mission
- Challenges
- Our solution
- Next steps


The AWS ML stack

Broadest and deepest set of capabilities















AI services

VISION			SPEECH		LANGUAGE		CHATBOTS	FORECASTING	RECOMMENDATIONS
 AMAZON REKOGNITION IMAGE	 AMAZON REKOGNITION VIDEO	 AMAZON TEXTRACT	 AMAZON POLLY	 AMAZON TRANSCRIBE	 AMAZON TRANSLATE	 AMAZON COMPREHEND & COMPREHEND MEDICAL	 AMAZON LEX	 FORECAST	 AMAZON PERSONALIZE


ML services

 Amazon SageMaker	Ground Truth	Notebooks	Algorithms + Marketplace	Reinforcement learning	Training	Optimization	Deployment	Hosting


ML frameworks + infrastructure

FRAMEWORKS	INTERFACES	INFRASTRUCTURE								
 TensorFlow  mxnet  PYTORCH	 GLUON  K Keras	 EC2 P3 & P3DN	 EC2 G4 EC2 C5	 FPGAS	 DL CONTAINERS & AMIs	 AMAZON ELASTIC CONTAINER SERVICE	 ELASTIC KUBERNETES SERVICE	 AWS IoT GREENGRASS	 ELASTIC INFERENCE	 INFENTIA


Why machine learning on Kubernetes?



Composability



Portability



Scalability



Amazon EKS: run Kubernetes in cloud

Managed Kubernetes control plane, attach data plane

Native upstream Kubernetes experience


Platform for enterprises to run production-grade workloads

Integrates with additional AWS services

Which user are you?



ML practitioner



ML ops team member

ML practitioners would like to have...

- Infrastructure abstraction
- Sharing and collaboration
- End-to-end workflow (build, train, test, deploy)
- Experiment management
- Hyperparameter tuning/optimization

ML ops teams would need to implement...

- An end to end platform
- Infrastructure abstraction
- Authentication and authorization support (multi-tenant access)
- Resource and quota management

Introducing Kubeflow

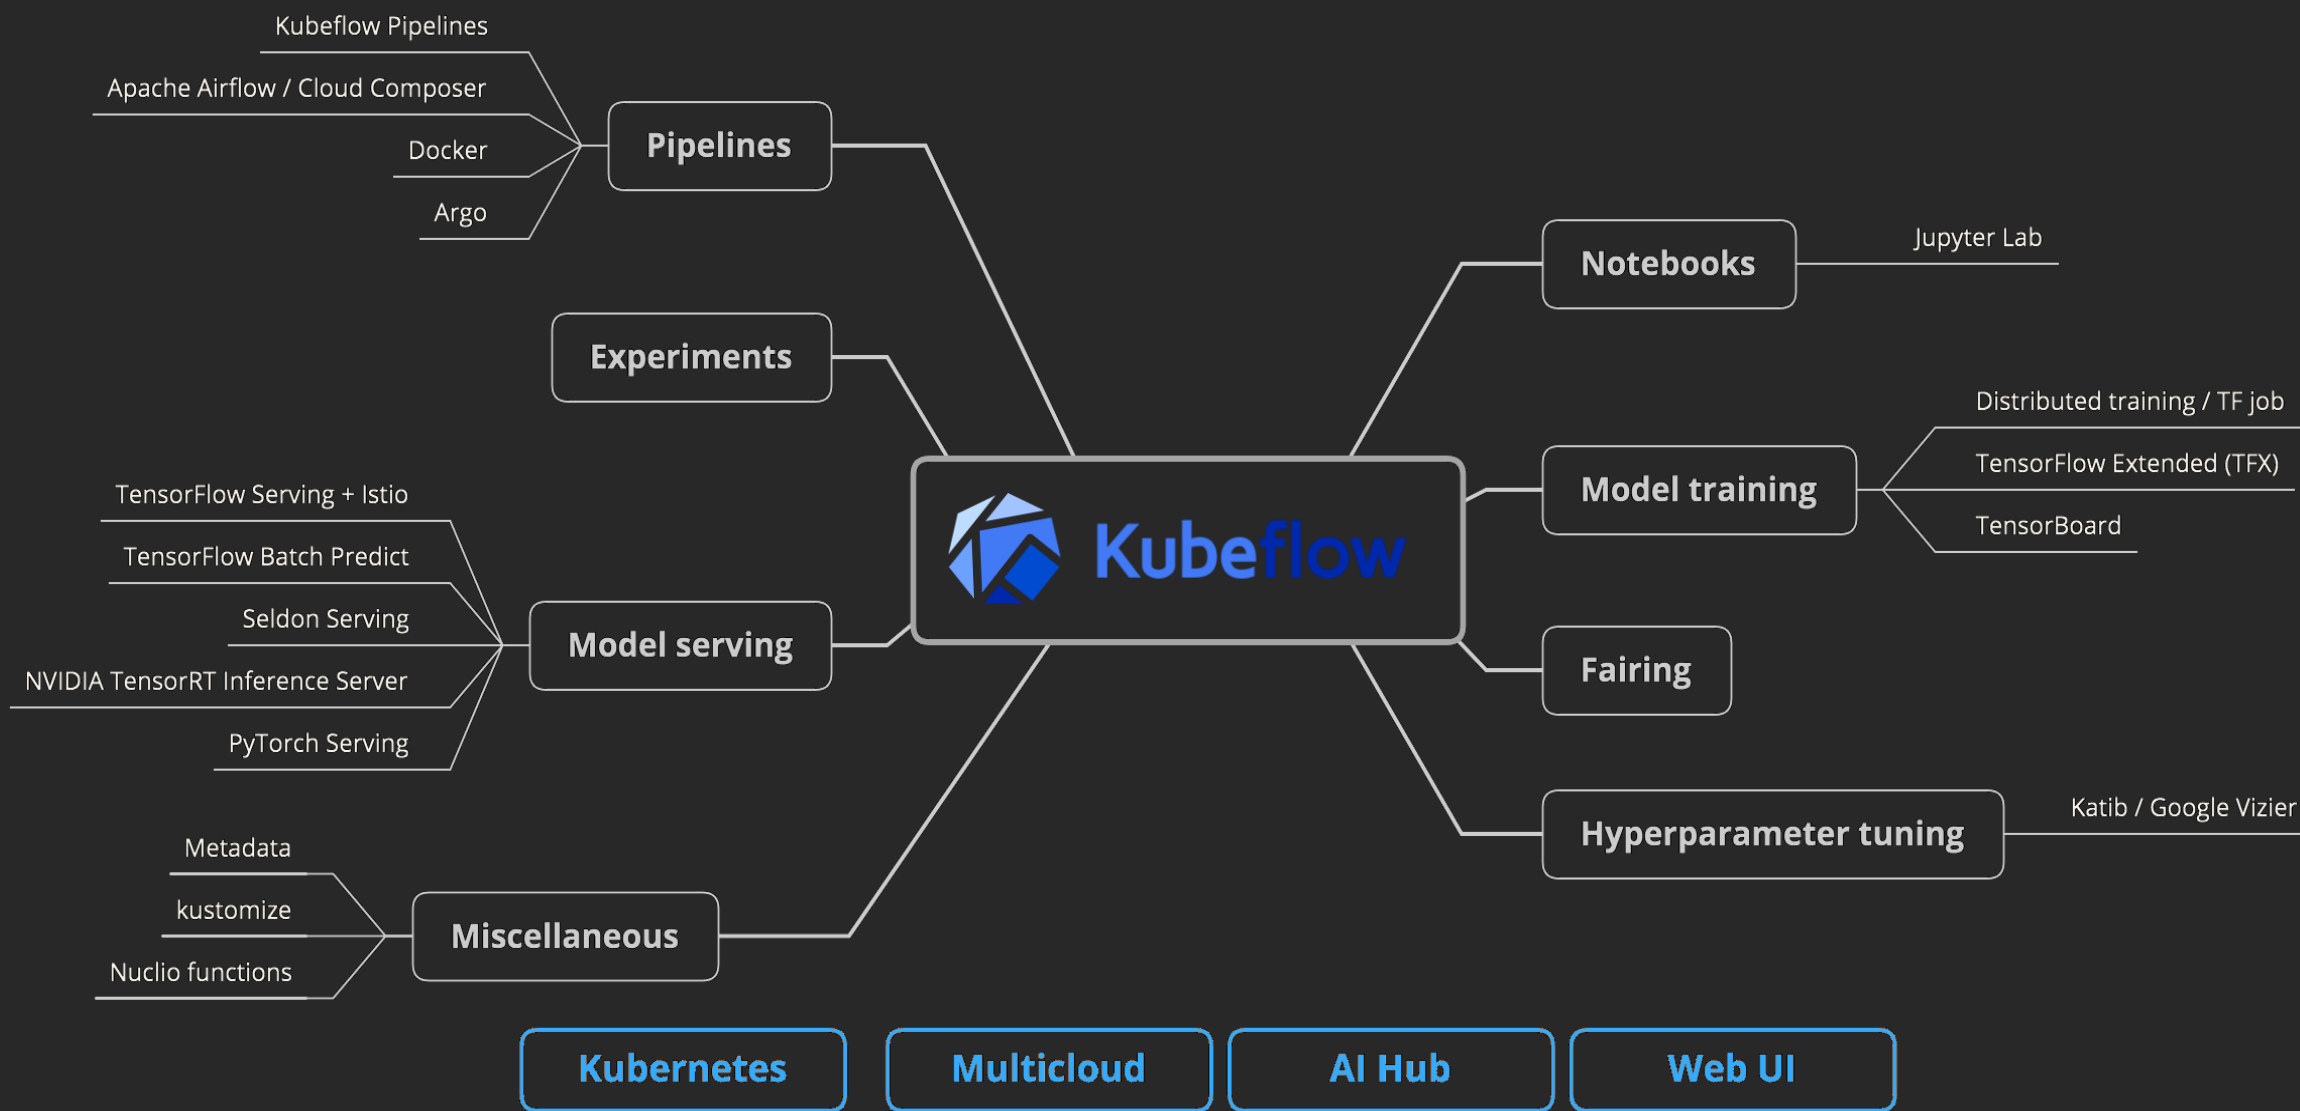


Notebook

Pipeline

Training

Serving



Prototyping

Jupyter / JupyterHub



- Build, deploy, and train ML models
- Live code, equations, visualizations, and narrative text
- 40+ programming languages
- Sharing and collaboration




EFS for reusing training data
and results

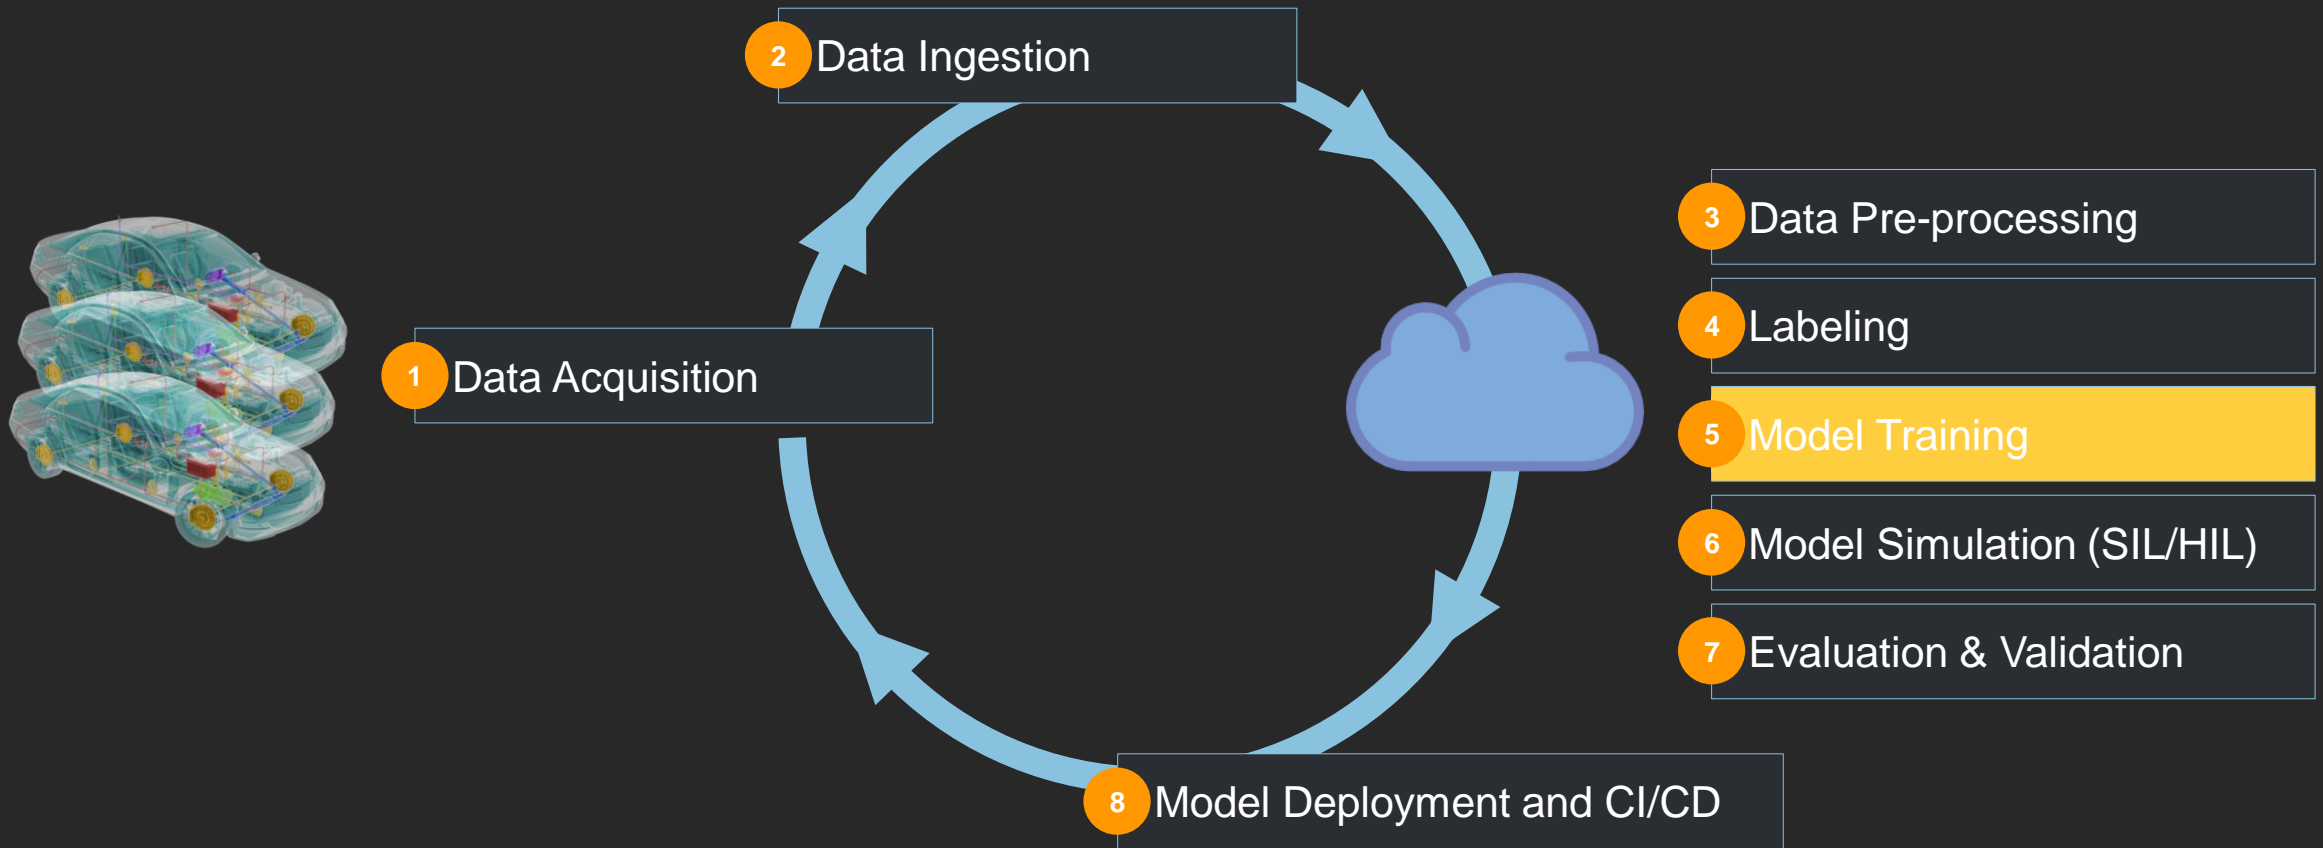


Built-in AWS CLI and ECR
support

Training



Typical Autonomous Vehicle Development Workflow



Distributed Training Challenges

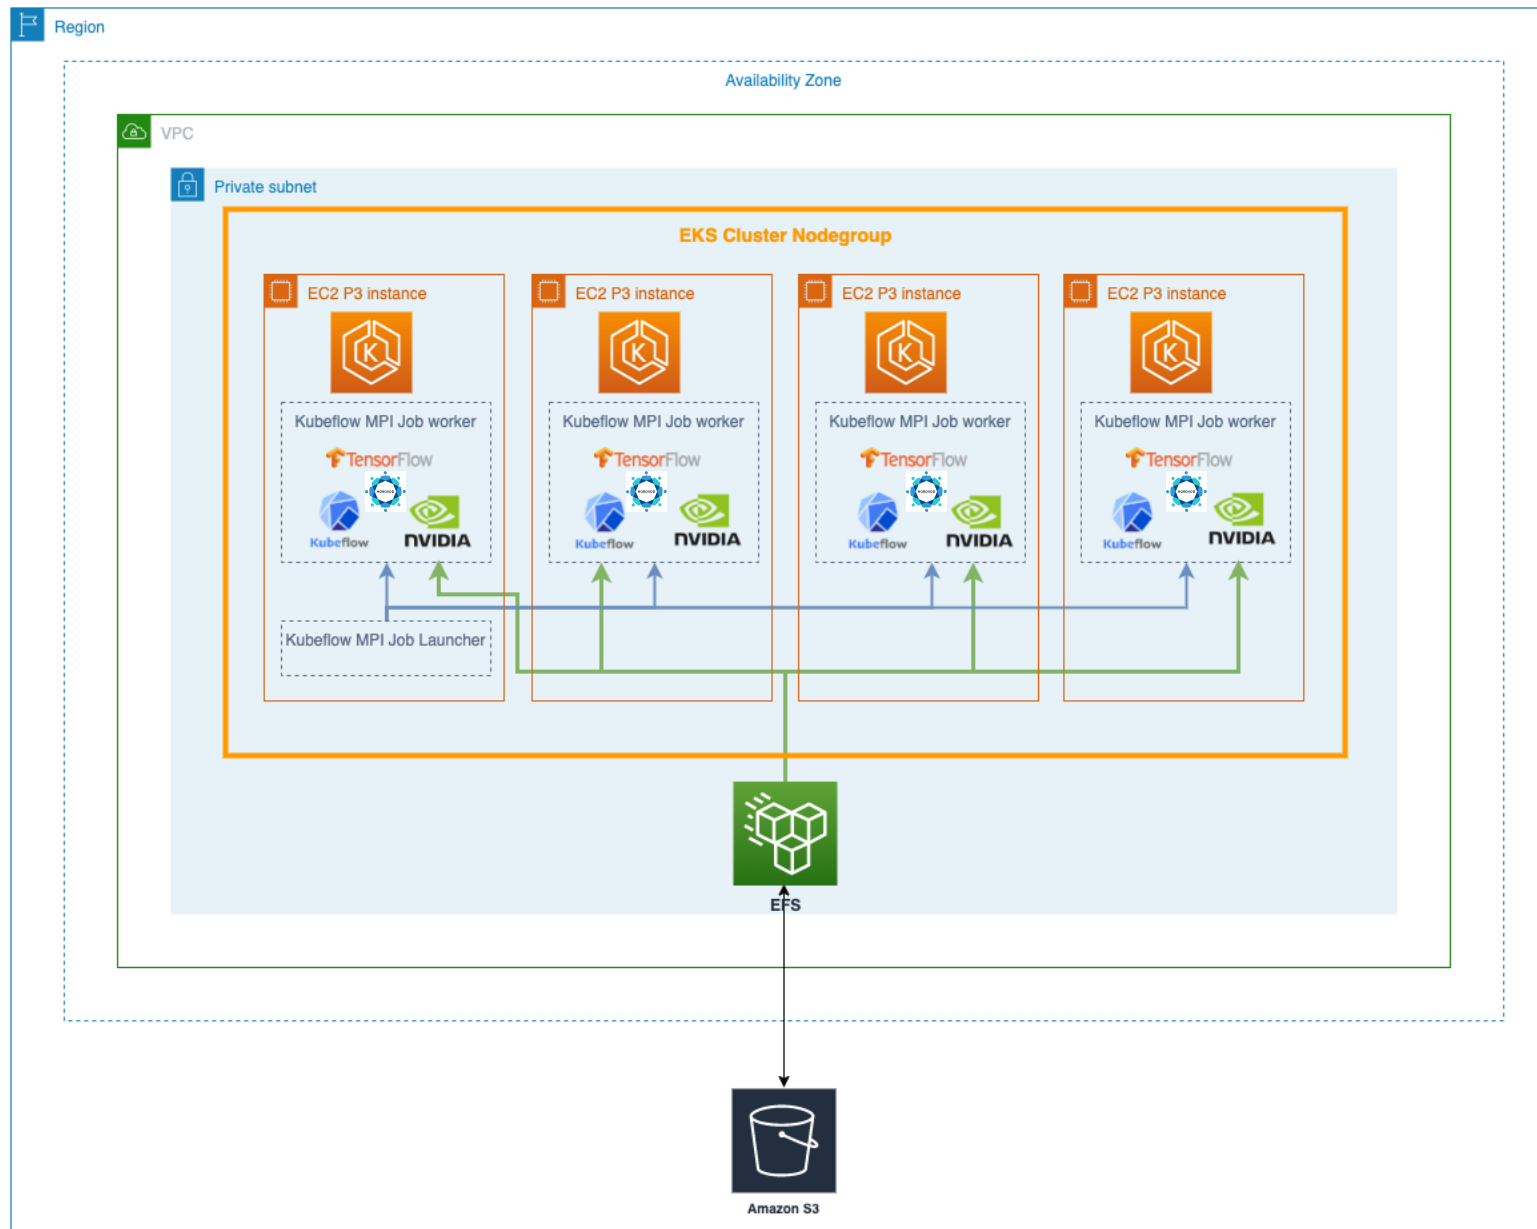
- Single GPU code → multiple
- Dataset Copying time
- Dataset Sharing and Reuse

Horovod + MPIJob

Use FSx Lustre / EFS



Built-in CSI driver with S3
integration



Distributed Training Challenges

- Single GPU code → multiple
- Dataset Copying time
- Dataset Sharing and Reuse

Horovod + MPIJob

Use FSx Lustre / EFS



Built-in CSI driver with S3
integration

Want to run Distributed Training on EKS?



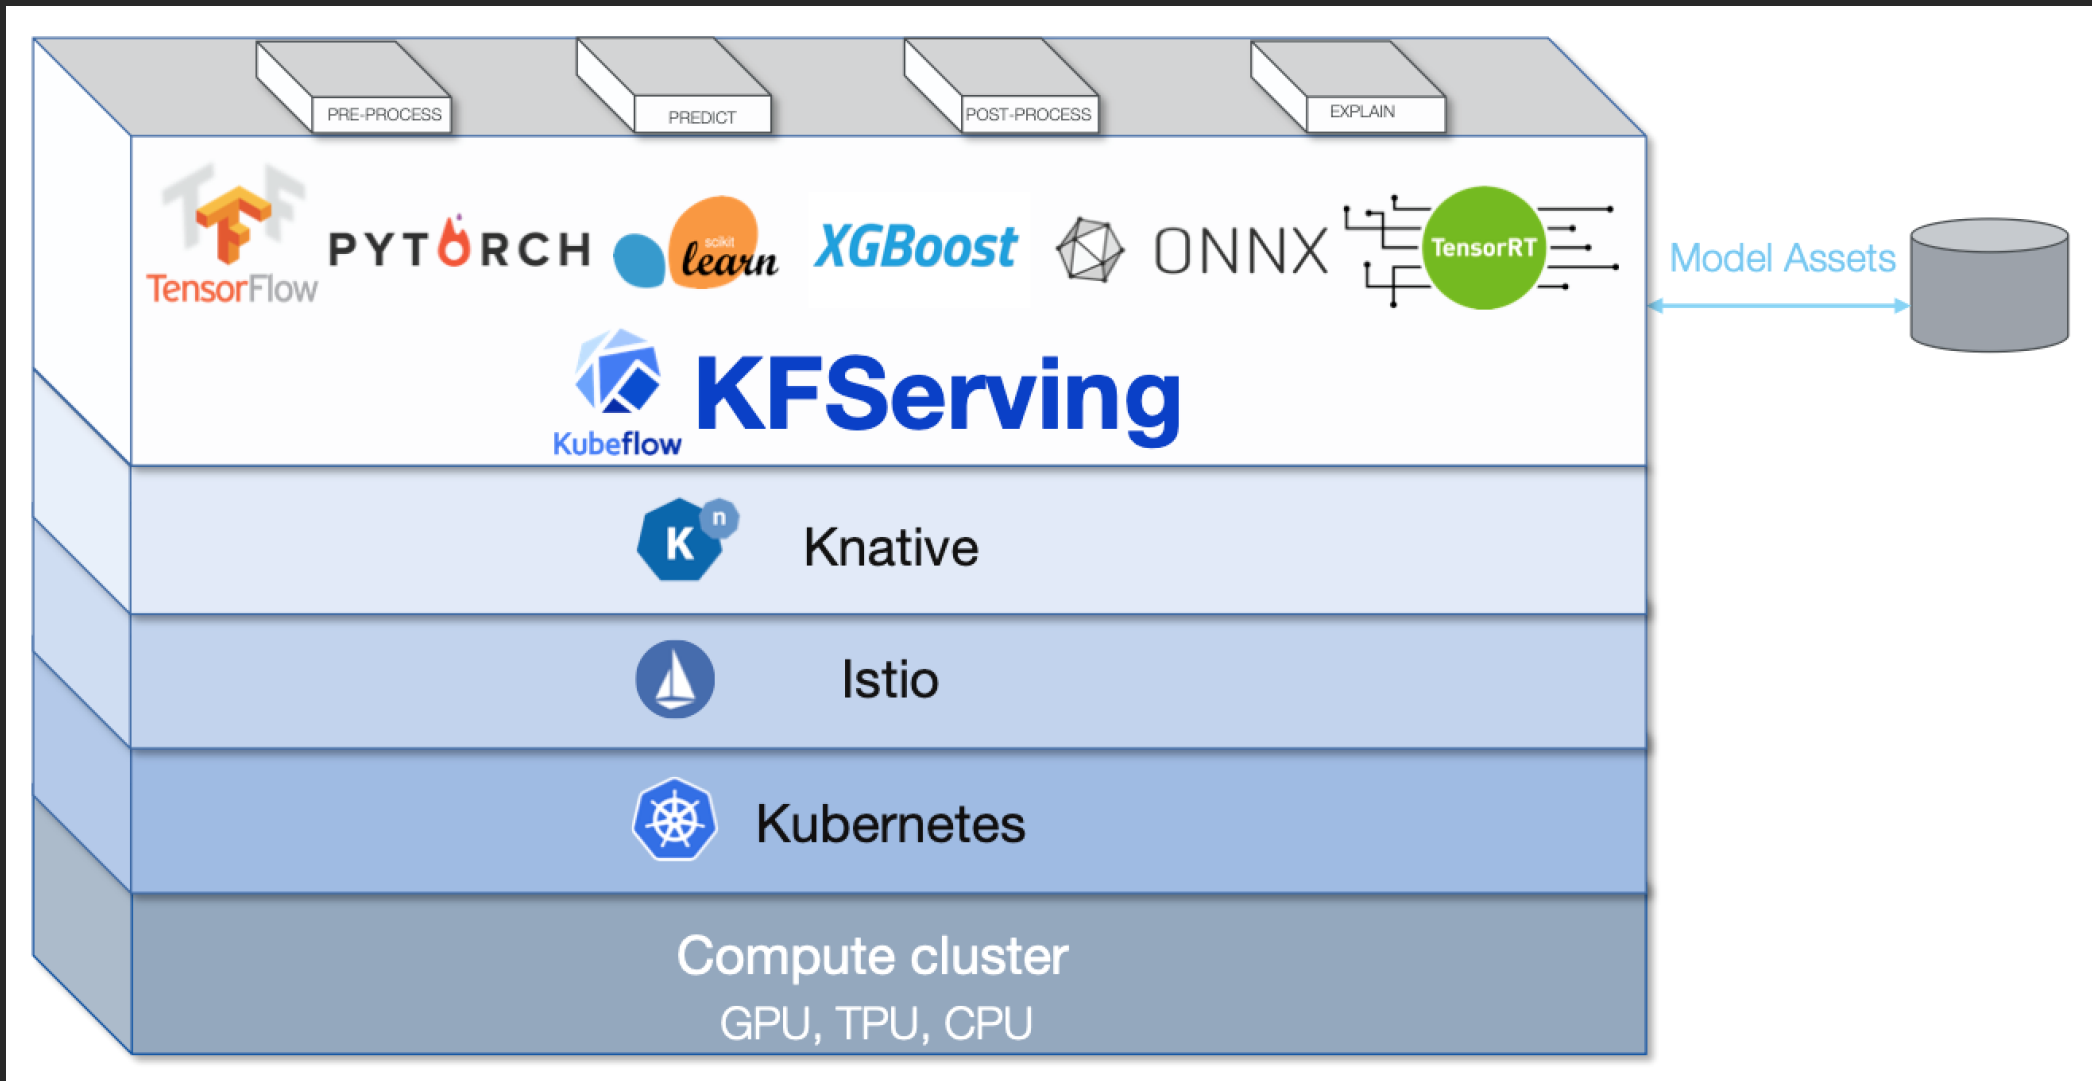
Distributed TensorFlow
training using Kubeflow on
Amazon EKS



Ajay Vohra
Principal SA -
Vision/AI/ML

Inference

Kubeflow KFServing



Pluggable Interface

```
apiVersion: "serving.kubeflow.org/v1alpha1"
kind: "InferenceService"
metadata:
  name: "sklearn-iris"
spec:
  default:
    sklearn:
      storageUri: "gs://kfserving-samples/models/sklearn/iris"
```



```
apiVersion: "serving.kubeflow.org/v1alpha1"
kind: "InferenceService"
metadata:
  name: "flowers-sample"
spec:
  default:
    tensorflow:
      storageUri: "gs://kfserving-samples/models/tensorflow/flowers"
```



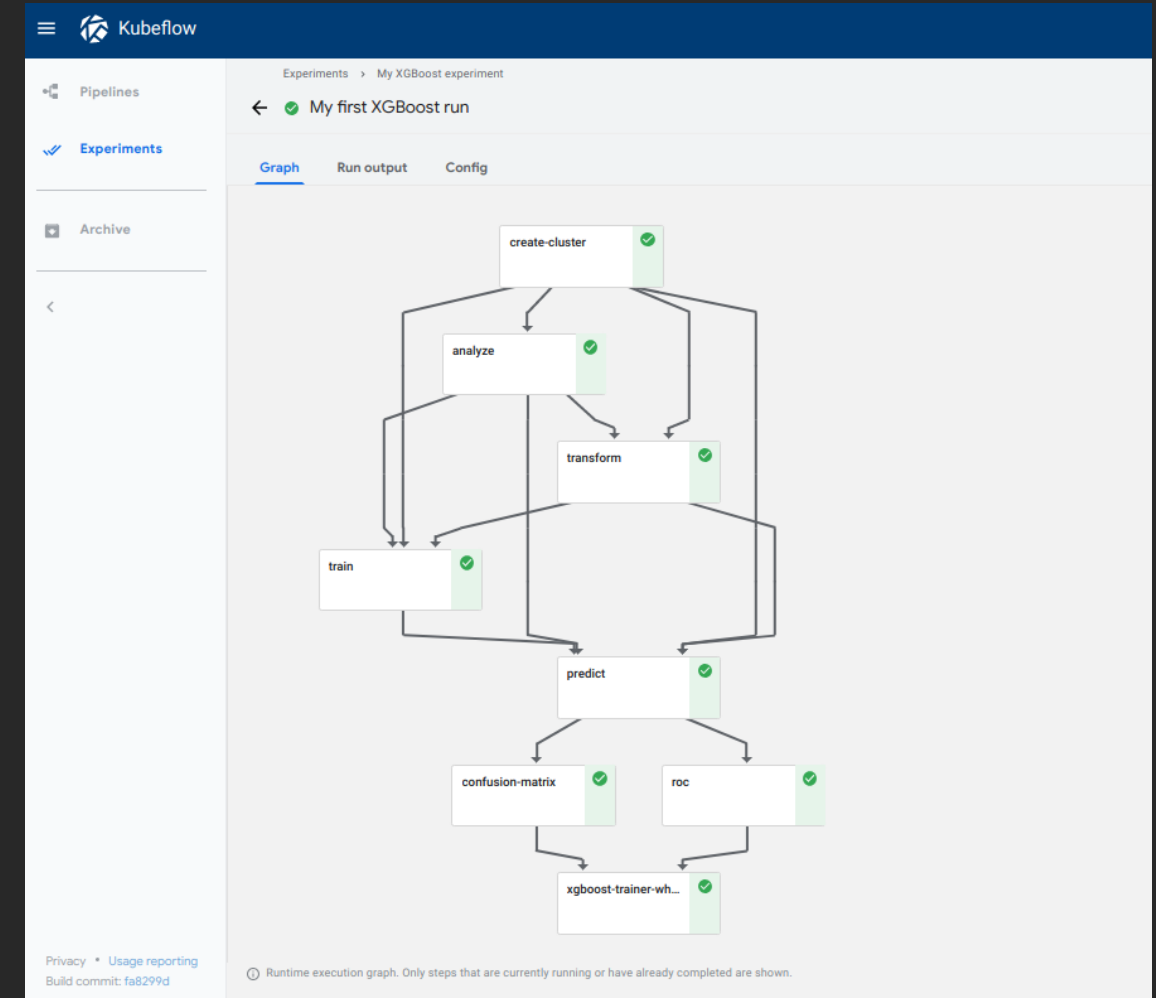
```
apiVersion: "serving.kubeflow.org/v1alpha1"
kind: "KFSERVICE"
metadata:
  name: "pytorch-cifar10"
spec:
  default:
    pytorch:
      storageUri: "gs://kfserving-samples/models/pytorch/cifar10"
      modelClassName: "Net"
```



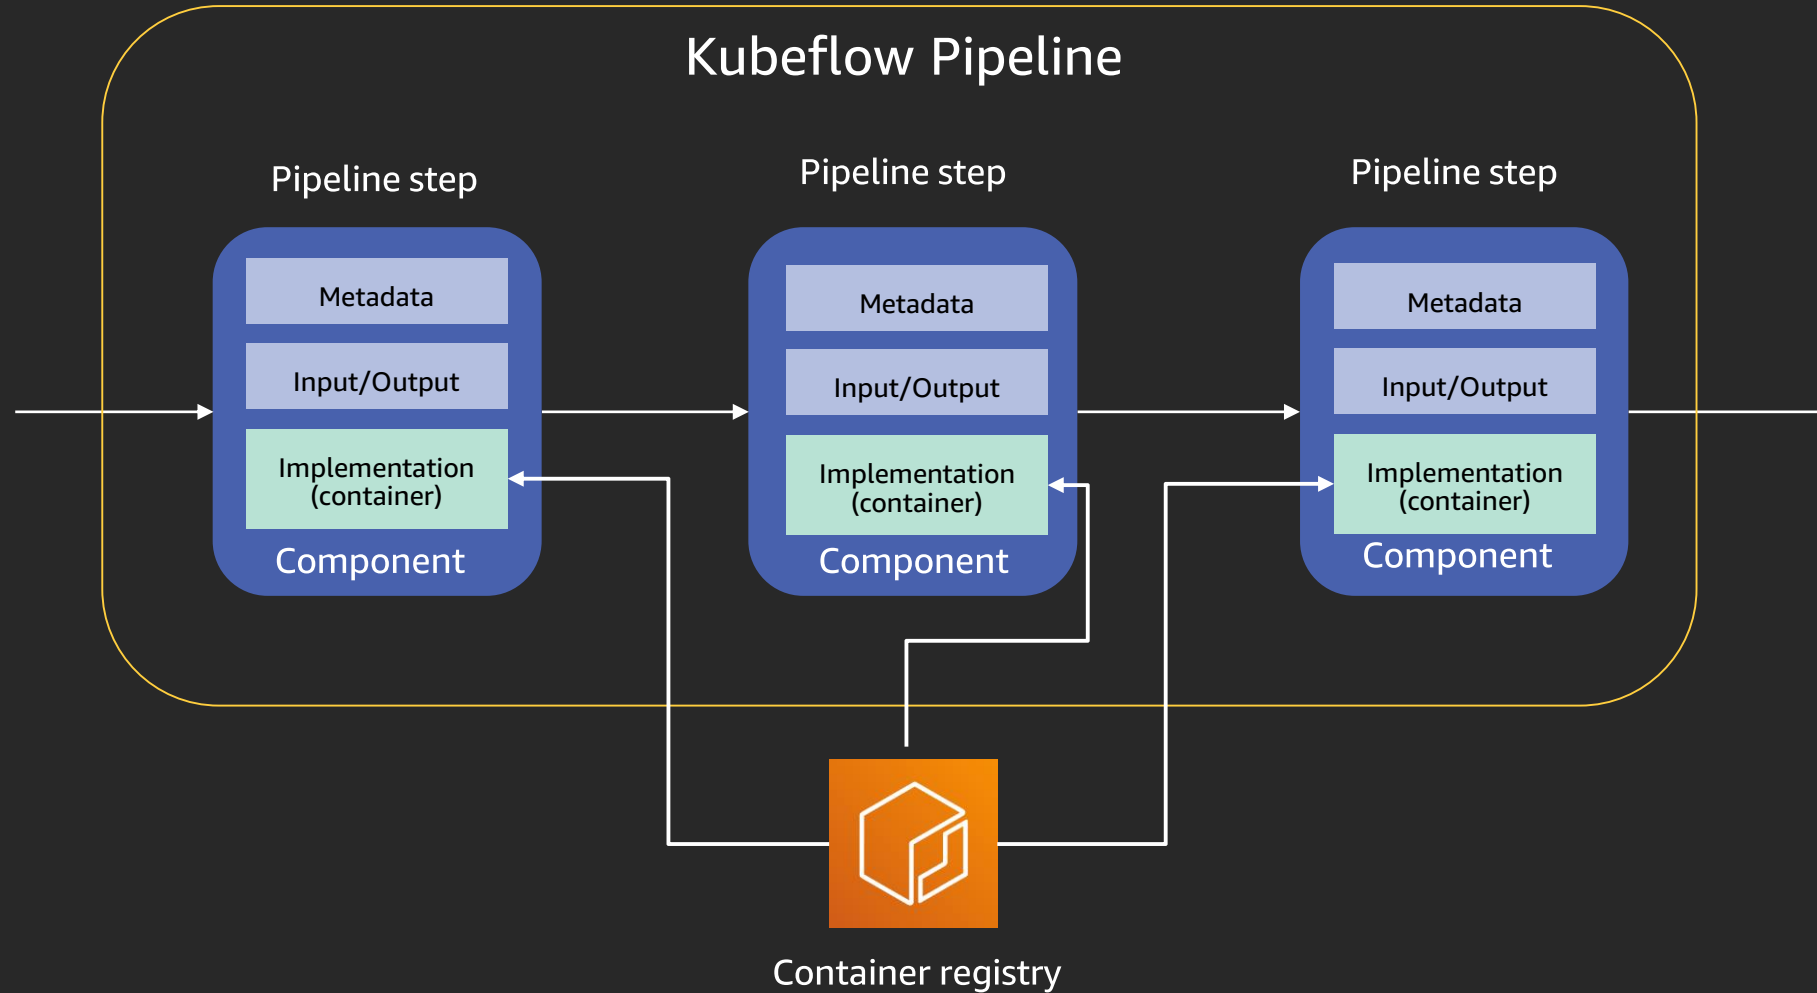
End-to-end ML

Kubeflow Pipelines

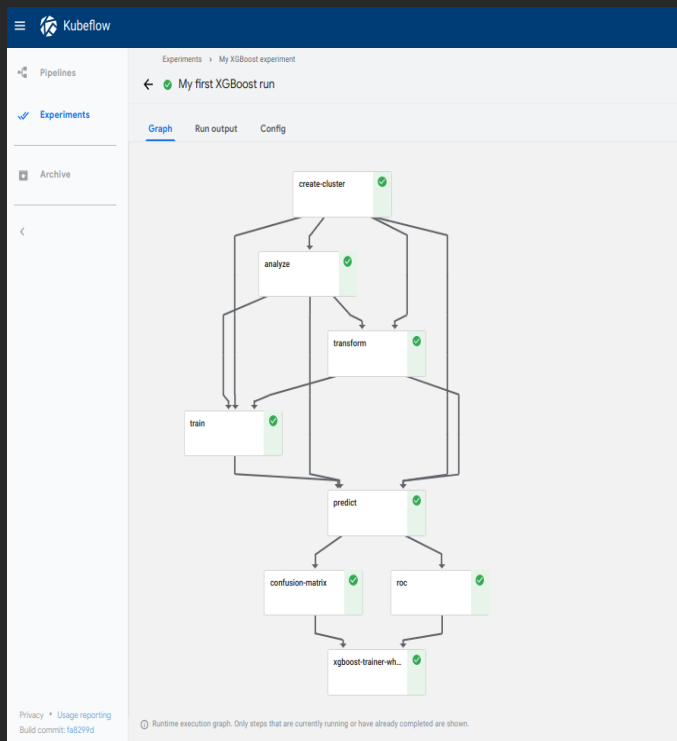
- A user interface (UI) for managing and tracking experiments, jobs, and runs
- An engine for scheduling multi-step ML workflows
- An SDK for defining and manipulating pipelines and components



Pipeline component



Creating a pipeline



Pipeline decorator

```
@dsl.pipeline(  
    name='Sample Trainer',  
    description=""  
)
```

Pipeline function

```
def sample_train_pipeline(...):
```

Pipeline component

```
    create_cluster_op = CreateClusterOp('create-cluster', ...)
```

```
    analyze_op = AnalyzeOp('analyze', ...)
```

```
    transform_op = TransformOp('transform', ...)
```

```
    train_op = TrainerOp('train', ...)
```

```
    predict_op = PredictOp('predict', ...)
```

```
    confusion_matrix_op = ConfusionMatrixOp('confusion-matrix', ...)
```

```
    roc_op = RocOp('roc', ...)
```



Compile pipeline

```
kfp.compiler.Compiler().compile(sample_train_pipeline, 'my-  
pipeline.zip')
```

Leveraging AWS Innovations through KubeFlow

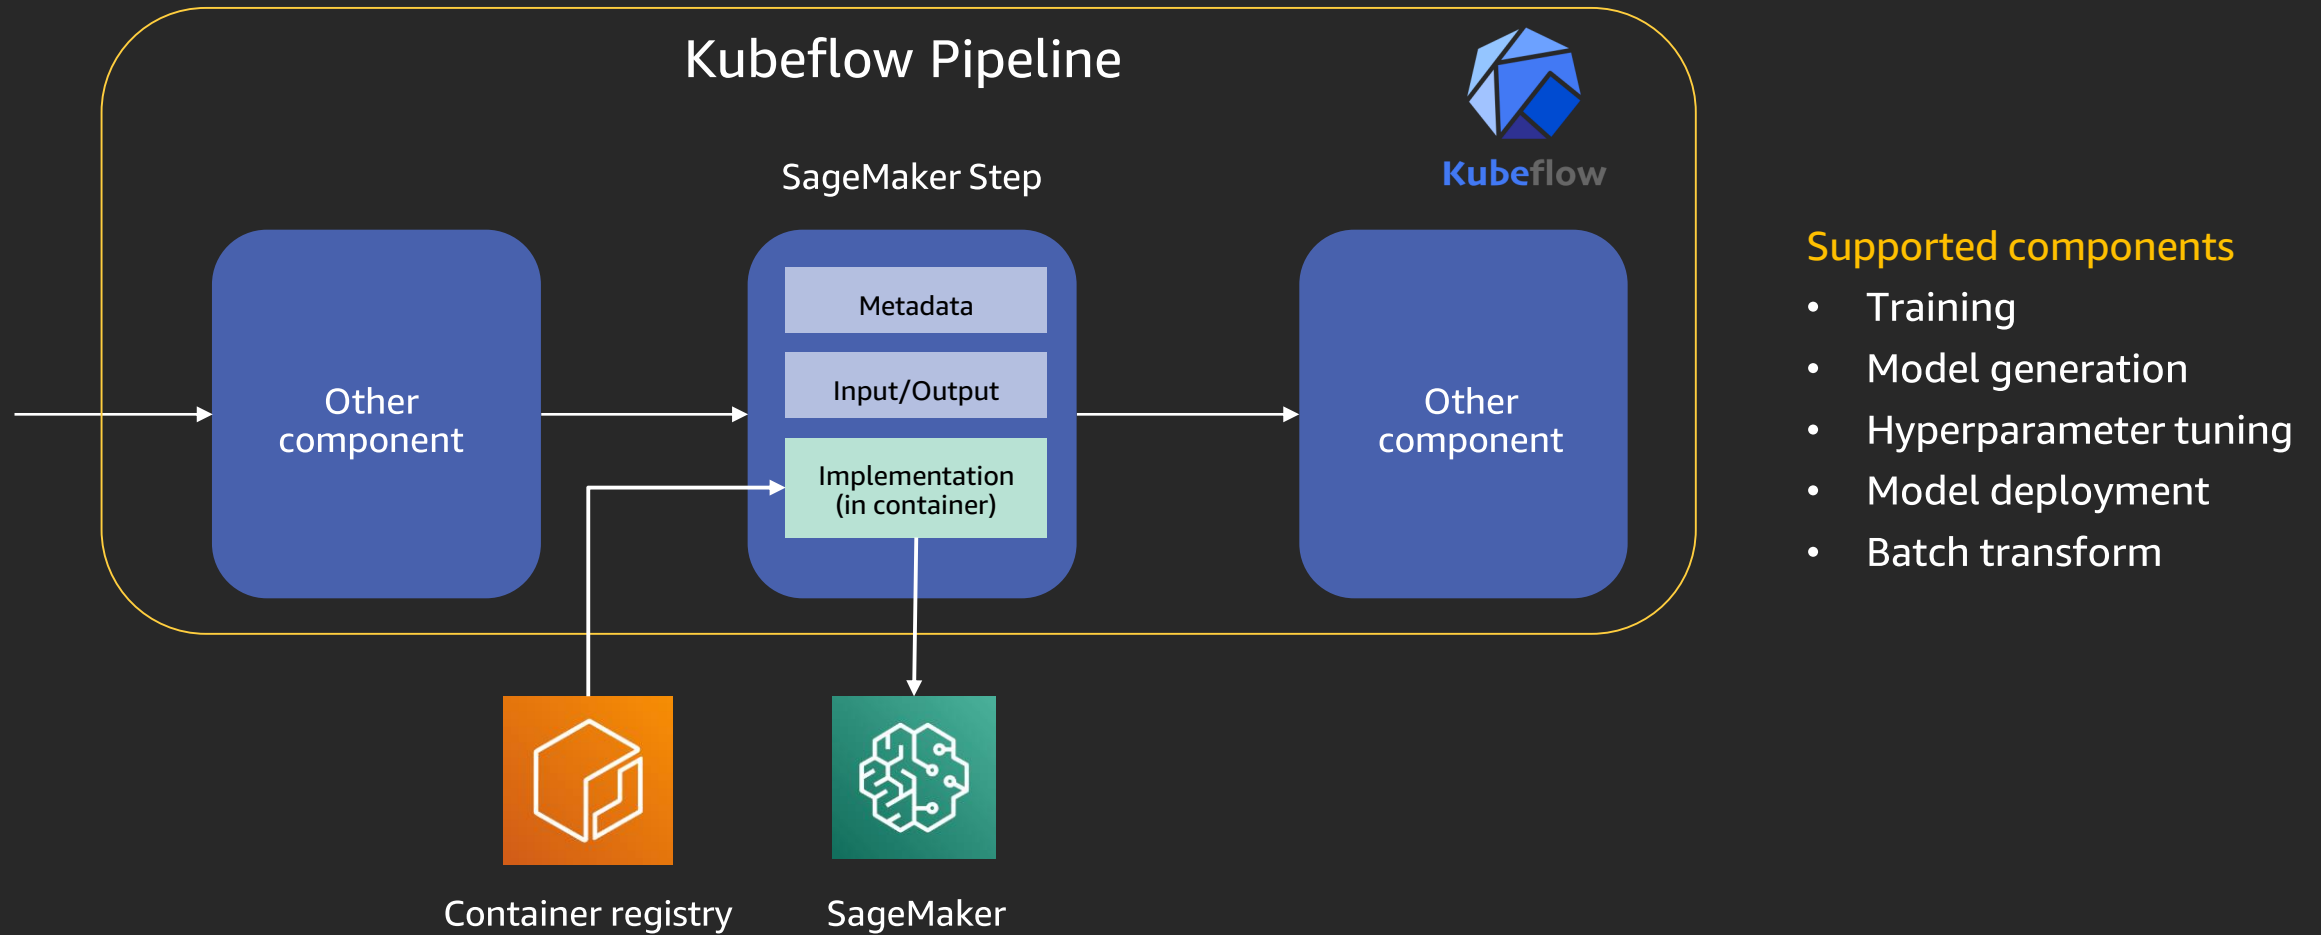


Do-It-Yourself



Managed Service

Running SageMaker – pipeline component



Github link <https://github.com/kubeflow/pipelines/tree/master/components/aws/sagemaker>

Related breakouts

[AIM326-R1] Implement ML workflows with Kubernetes and Amazon SageMaker

12/4/19 (Wednesday) 4:00 PM - Bellagio, Monet 1

What's next?

Kubeflow 1.0 – Main components

- **Graduating 1.0**
 - kfctl for deployment and upgrades
 - TFJob and PyTorch for distributed training (already 1.0)
 - Jupyter notebook controller and web app
 - ...
- **Beta**
 - Katib for hyper-parameter tuning
 - Fairing SDK to facilitate use of notebooks for build-train-deploy
 - KFServing for model deployment and inference
 - ...

Kubeflow 1.0 – AWS Support

- Multi user support
 - Kubeflow pipelines
 - Managed contributors
- IAM Roles for Service Accounts integration with notebooks

Call to Action

re:Invent workshop

OPN401-R1 - [REPEAT 1] Machine learning with Kubeflow on AWS

12/5/19 (Thursday) 3:15 PM - MGM, Level 1, Grand Ballroom 120

Online workshop

<https://eksworkshop.com/kubeflow/>

Community

Join the **kubeflow#aws** Slack channel:



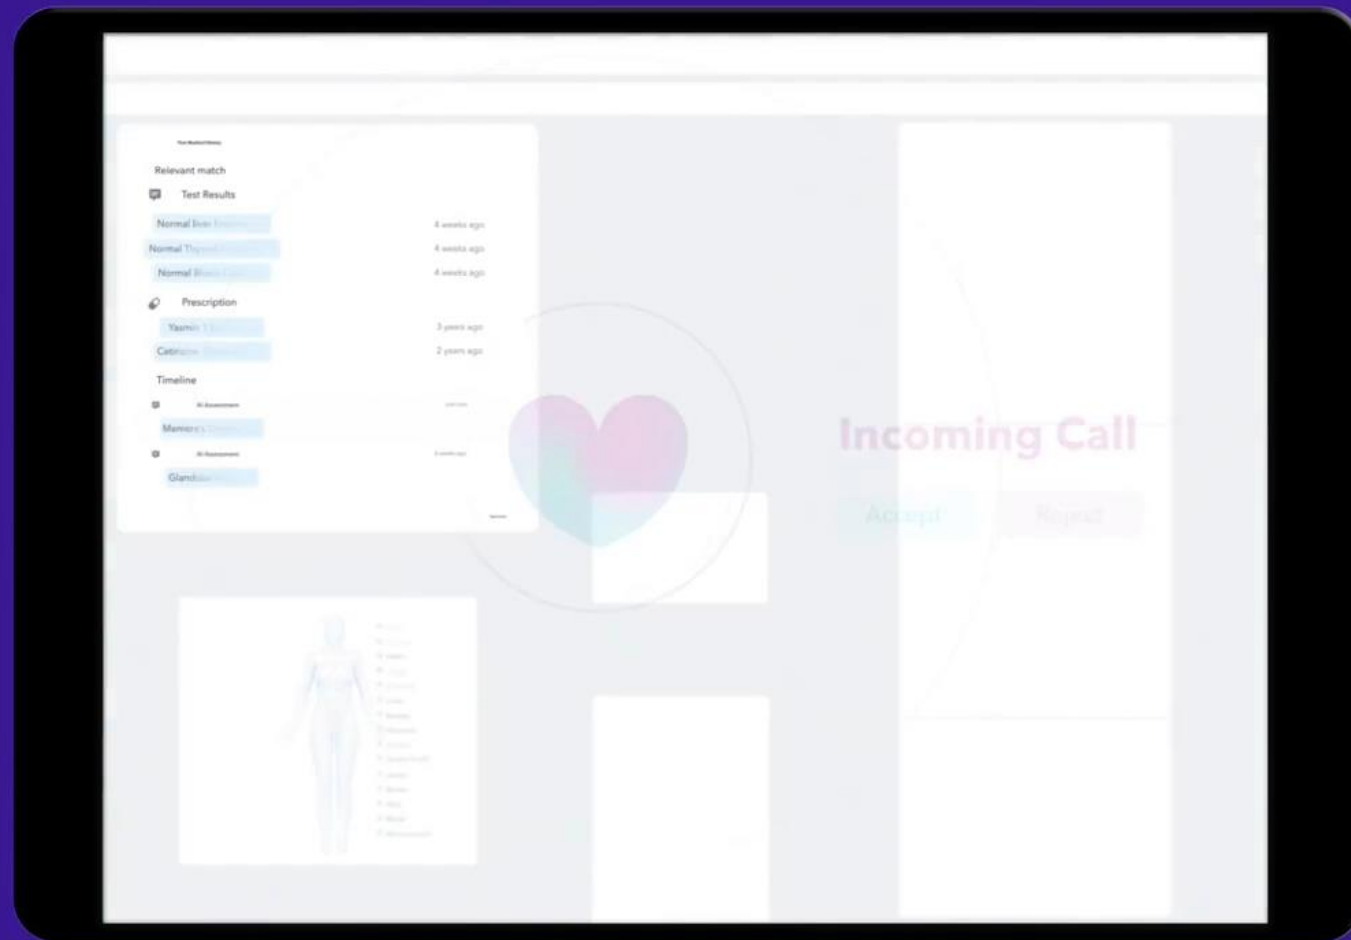
Babylon

Babylon

The mission



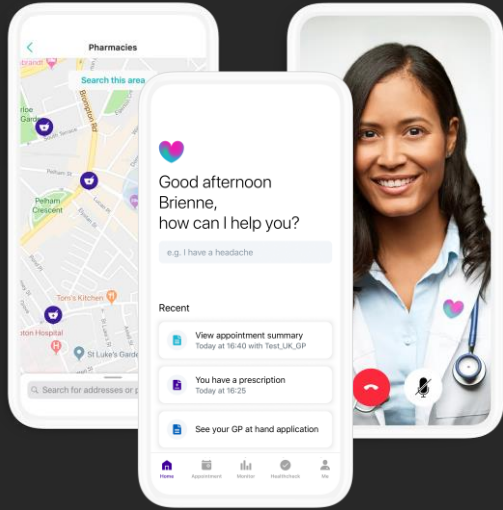
We believe it is possible to put
an **accessible** and **affordable**
health service in the hands of
every person on Earth



Babylon

Challenges

Our challenges



Global products



Global company



Global platform

Our challenges

300+

Microservices

500+

Deployments per
day on average

15+

EKS clusters

5

AWS Regions

Our challenges

- Providing a **safe** and **secure** environment for our research teams
- Data locality
- Training and managing AI models at a **global scale**
- Improving overall engineering **efficiency**

Babylon

Our solution

AI platform: The big picture


Providing our teams with a single interface to **experiment**, **train**, **tune**, **validate**, and **track** their AI models

AI platform



Secure

Treat it as a
production
platform



Scalable

Pay for what
you use, and
scale as much
as you need



Flexible


It must be able
to run any tool
or service our
users need



Global

We need to
deploy it
anywhere we
operate

AI platform



AI platform



- EKS infrastructure
- GPUs
- Networking




- Kubeflow deployment
- Project structure
- Extensibility



- Global platform
- Deployments

AI platform: *EKS infrastructure*

- Make everything private
 - EKS API
 - EKS nodes
- Prepare for different workloads
 - High CPU, high memory, GPUs
 - Spot instances for better pricing
- Encrypt by default
 - EBS volumes
 - AMIs (root volumes)



AI platform: *GPUs on EKS*

- NVidia device plugin *DaemonSet*
- Register nodes with
 - Taint: `nvidia.com/gpu=true`
 - Label: GPU type and brand
`nvidia-tesla-k80`

```
# Add a toleration, allowing the pod to run on GPU nodes
tolerations:
  - key: "nvidia.com/gpu"
    operator: "Equal"
    value: "true"
    effect: "NoSchedule"

# Use nodeSelector to choose the GPU type
nodeSelector:
  accelerator: nvidia-tesla-k80 # or nvidia-tesla-v100
```

AI platform: *Networking*

- Virtual Private Cloud (VPC)
 - Span over at least 3 Availability Zones (AZ) ... if you can!
 - 1 *Auto Scaling Group* per AZ per instance type
 - Use NAT gateways
- Kubernetes
 - Zero-trust policy
 - Mutual TLS
 - Service role-based access control
 - Service mesh (Istio)

AI platform: *Kubernetes and Kubeflow*

- Kubernetes


- Perfect for scheduling and running any workloads
- Massive scale

- Kubeflow

- ML toolkit for Kubernetes
- Modular: use and deploy only what you need
- Multi-user: plug in your enterprise OIDC
- Open-source: deploy it anywhere you have Kubernetes
- We mostly use Jupyter Notebooks, TensorFlow jobs, hyperparameter tuning

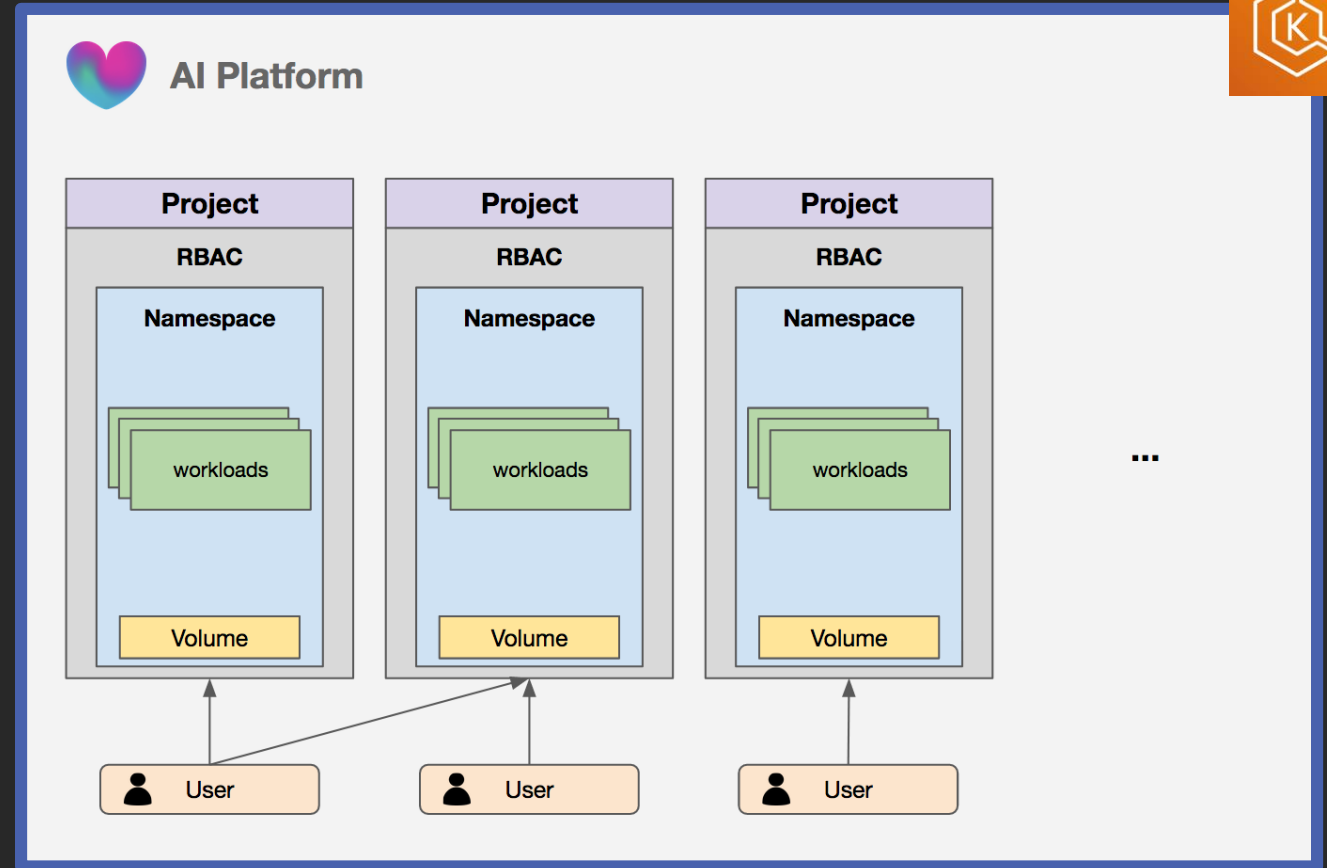


AI platform: Deploying *Kubeflow*



AI platform: *Project isolation and collaboration*

- Project *CRD*
 - Isolated namespace
 - RBAC rules for user management
 - Encrypted **EFS** partition
 - Quotas management
- Need faster storage?
 - EBS volumes
 - FSx for Lustre



AI Platform: *Workload Monitoring*

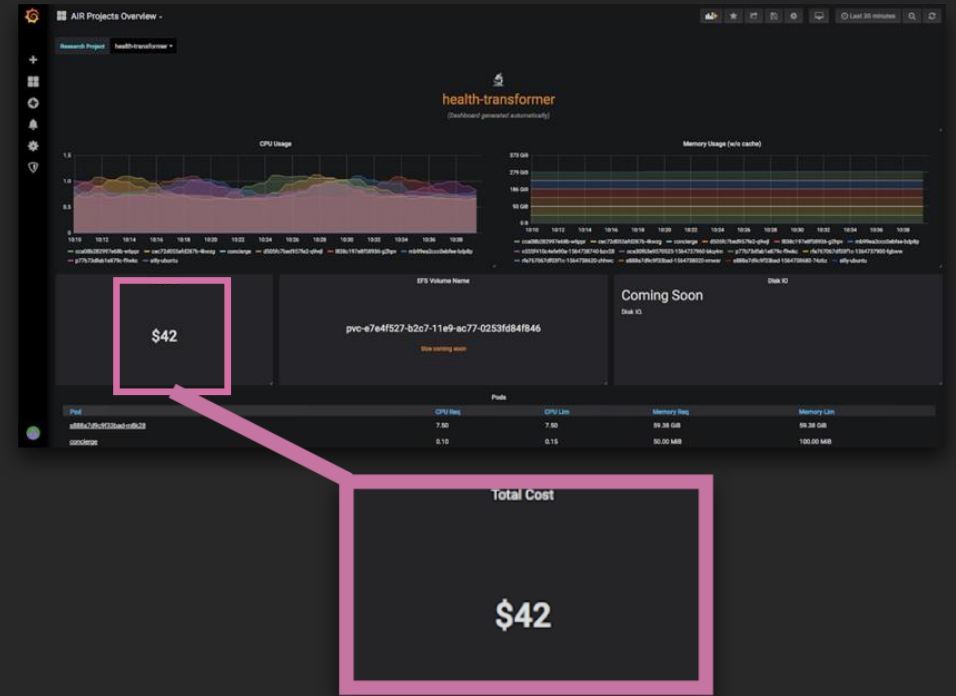
- Monitor Projects (*CPU, RAM, etc*)
- Monitor Cluster State (*MLOps*)
- Collect metrics from jobs
- Automated Dashboards via ConfigMaps



Prometheus



Grafana



... And monitor cost

AI platform: *Extensibility of services*

- Modularity of Kubernetes and Kubeflow

- Gives us ability to add new tools *fast*
- Ex: deploying *Argo* for workflow management

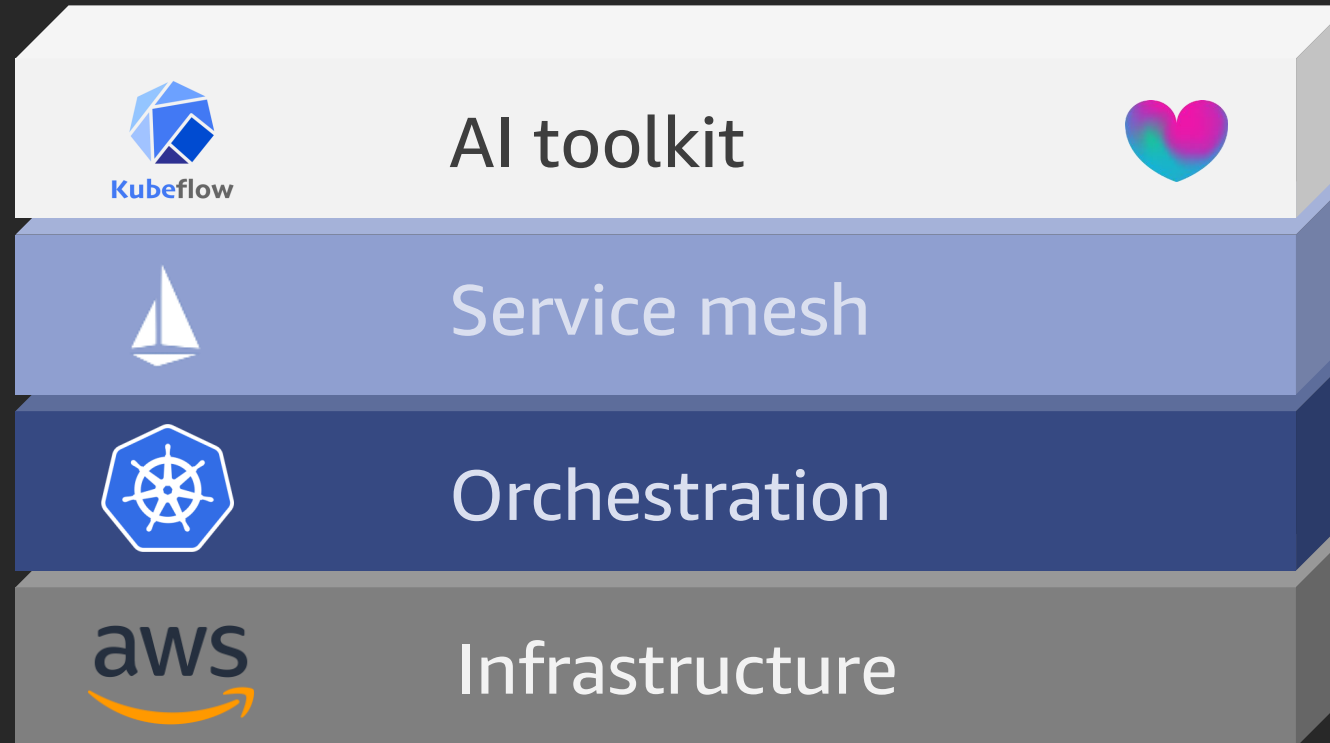


- Use case:

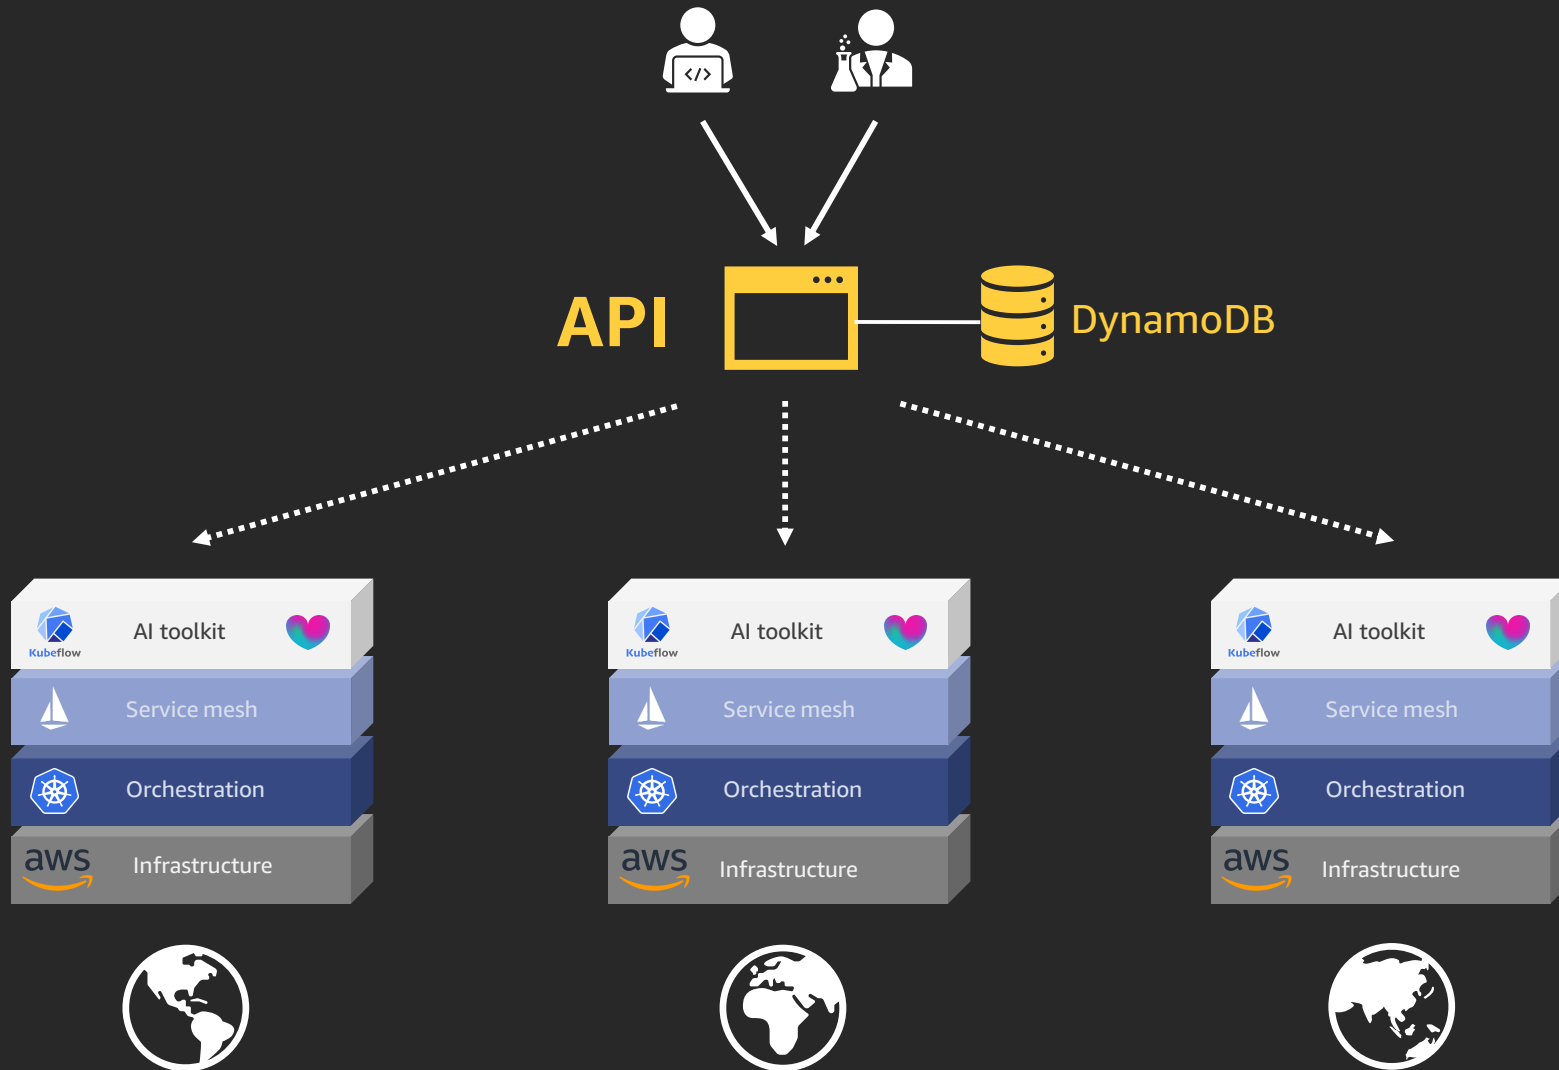
- Migrating our model Clinical Validation pipeline to Argo on EKS



AI platform: *Single* cluster



AI platform: *Global footprint*



API



```
{
  "kind": "pod",
  "name": "simple-gpu-example",
  "image": "nvidia/cuda:8.0-cudnn5-runtime",
  "command": ["python"],
  "args": ["script.py"],
  "resources": "gpu_medium"
}
```



```
---
apiVersion: v1
kind: Pod
metadata:
  name: simple-gpu-example
  namespace: my-project
spec:
  containers:
  - image: nvidia/cuda:8.0-cudnn5-runtime
    command: [ "python" ]
    args: [ "script.py" ]
    name: simple-gpu-example
    resources:
      limits:
        memory: "16Gi"
        cpu: "8000m"
        nvidia.com/gpu: 1
    name: "tensorflow"
    volumeMounts:
    - mountPath: /mnt
      name: efs-storage
  restartPolicy: "OnFailure"
  volumes:
  - name: efs-storage
    persistentVolumeClaim:
      claimName: efs
  imagePullSecrets:
  - name: my-deploy-pull-secret
  tolerations:
  - key: "nvidia.com/gpu"
    operator: "Equal"
    value: "true"
    effect: "NoSchedule"
  nodeSelector:
    accelerator: nvidia-tesla-k80
```

Babylon platform: *Global deployments*

- Deploying models and services in multiple regions
- Secure SDLC-aware deployment tool: *Shipcat* (in Rust)
- Per microservice tracking of:
 - Compliance and regulations
 - Engineering documents
 - Data management



Babylon

Next steps

Next steps

- Providing a common **model serving** framework at Babylon
- Better **metadata tracking** for our models
- Integration with **Kubeflow pipelines**
- Improving Docker **user experience** for researchers

Thank you!

Yaniv Donenfeld

donenfel@amazon.com

Jean-Marie Ferdegue

jeanmarie.ferdegue@babylon
health.com

Jeremie Vallee

jeremie.vallee@babylon
health.com



Please complete the session
survey in the mobile app.